

September 15, 2020

1 Individual Assignment 2

1.0.1 (Due on Sep 8 11:59PM on LumiNUS)

Student Name: Haojun Gao

Student ID: A0218869J

1.1 Problem 1 (GrabWheels)

Towards the smart nation initiative, the Government of Singapore has encouraged sustainable and smart mobility solutions such as recent scooter vehicle sharing (e.g., GrabWheels). Launched in the late 2018, GrabWheels rolled out around 30 parking locations in NUS Kent Ridge campus. For more details, please see <http://news.nus.edu.sg/press-releases/new-e-scooter-sharing-service> and <https://www.grab.com/sg/wheels/>. A group of DBA5106 students is working on a project in the design and operations of GrabWheels at NUS campus, e.g., where to set up the parking locations and how many scooters to place at each location. Please answer the following questions.

a) Suppose GrabWheels wants to improve its service quality by optimizing parking locations and calculating the optimal number of scooters at each location. Please suggest what information would help and what data GrabWheels needs to be collected. [2 pts] There are two ways to improve the service quality: - optimize the infrastructure location (optimizing parking locations) - solve the relocation optimization problem (cope with the problem caused by demand imbalance)

Therefore, from the data type point of view, both spatial data and time series time will be helpful. And if we look at the data source, shared scooter vehicle usage data (taking the vehicle as a unit can avoid the use of user data and reduce the collection of user data) and other sort of data that can affect user behavior could be useful. - Infrastructure Information - Parking locations - Number of trips that were taken each location each day - Scooter Vehicle Trip Dataset - vehicle ID - trip starting time - trip end time - trip start location - trip end location - Weather Dataset - Temperature: basic indicator - Precipitation: an indicator of whether rain or not - Wind speed: an indicator of whether suitable for taking scooter - Date-related Dataset - holiday - etc. - Spatial Dataset - Topographic features of parking locations - Distance from the nearest parking location - etc.

b) By Sep 2019, GrabWheels has collected 2-month data (you specified in part a) from its operating 30 locations. To forecast the demand for new locations, e.g., COM2, briefly discuss which data mining task should be performed and what additional data sets should be collected, if any. [2 pts] Data Mining Task

- We could define the problem as a regression problem.
- Forecast objective is the average daily demand of a specific parking location. We could use the 2-month data mentioned above to form the variable(average demand of the parking location). And use spatial data as features to train the model.
- The collected data could be divided into train, validate and test set to model and evaluate.

Additional Data Sets

- As for the additional data sets, the relevant feature data of those new locations await predicted also need to be collected.

c) Following part (b), briefly discuss how GrabWheels can forecast the demand scooter sharing at new locations, using your proposed approach. Please identify which steps are the process of data mining (DM), or the use of the results of data mining (Use). [2 pts] According to the model proposed in the previous question, we basically use the geographical characteristics to predict the use demand of this location. The basic idea is that the inherent characteristics of a place will determine, to a great extent, the demand of the scooter in this place.

For example:

- if this place is close to the bus station or MRT, the demand for scooter may be great.
- If this location is close to the giant company, the demand for scooter may also be great.
- However, if the other/nearest parking location is far away, then one may not use scooter as the transportation.

Process of Data Mining

- Find the feasible locations as a new parking locations.
- Collect the geographical features of existing parking locations and feasible locations.
- Use the existing location data to train and evaluate the model.
- Predict the average demand of the new locations

Knowledge Discovery and Decision Making

- The prediction results could be used as evidence to support or influence the company's decision.
- The feature importance could also perform as an application for knowledge discovery, such like which feature is most relevant to the use of scooters.

d) Discuss what challenges GrabWheels may face and how data analytics can help improve its operations. [2 pts] There are still many tough situations that challenge GrabWheels and also threaten their services, and demand imbalance is the most critical one among them. To a large extent, bike sharing systems are used for one-way trips, and such a trend leads to inappropriate bike distribution in time and space. Consider the scenario: - Many people choose to ride to this place using the scooter, but few rides away from this place. That will cause the scooter to pile up, and the utilization rate drops sharply. - Conversely, if many want to ride away meanwhile few people ride over here, this condition also results in an unsatisfactory service and leads to the low revenue.

To tackle this dilemma, the redistribution of bikes over stations is required. The straightforward strategy is to predict the real-time demand(hourly), rather than the average need of the parking

location. Then turn the problem into a path optimization problem: how to relocate the scooter using the least human resources and finance cost.

- The first step to optimize the relocation of scooters in the system is to be able to predict the number of scooters in a particular parking location at particular time, which can be done through the use of machine learning algorithms.

e) If you are evaluating the market expansion plan of GrabWheels, e.g., whether it should cover certain regions, describe what data you would like to collect and how you would like to collect and analyze them for such an evaluation? [2 pts] When it comes to making an expansion plan, model could be level up to the region level, instead of location level. It consists on predicting the total demand of scooters in the region, both returning and renting. This model could provide general information on the number of scooter demand in a particular zones, where it is not necessary to know exactly the decision of the parking location selection.

Region level information is needed to construct the model. Such like, region area, prosperity of the region, region population, age composition of the population, etc.

Same as the above model, we could use the operation performance of the known regions to train the model and predict the performance of the unknown region. And we could use the model results to decide which region could be a target region that needs to be covered in the future.

1.2 Problem 2 (Blue Bikes)

A renowned consulting firm MSBA & Company is currently analyzing the trip data of Blue Bikes (originally Hubway) in Boston. Blue Bikes (<https://www.bluebikes.com/>) is a public bike share system that operates in the Greater Boston area.

We are interested in exploring the bike share operations from the data available on <https://www.bluebikes.com/system-data>. We will utilize the “Bluebikes trip history data” as well as “the list of GPS coordinates and number of docks for each station” available on the website.

Please download and analyze the data to answer the following questions. [Note: (1) the unzipped “201906-bluebikes-tripdata.zip” has a wrong file name (with correct data inside); (2) you can add extension “.csv” to the unzipped “201907-bluebikes-tripdata” file]

1. Provide the line chart of monthly trips from 2015-06 to 2020-05. [2 pts]

```
[1]: import sys
import matplotlib
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: sys.path.append(
    r'D:\OneDrive\Programming\documents_python\NUS_
    ↳Courses\DBA5106\Course-DBA5106\assignment\IndividualAssignment2')
```

```
[3]: from utils.logger import logger
from utils.config import PROCESSED_DATA_DIR
from utils.data_porter import read_from_csv
```

```
[4]: tripdata_df = read_from_csv('tripdata.csv', PROCESSED_DATA_DIR,
                                parse_dates=['starttime', 'stoptime'])
```

```
[5]: tripdata_df.dtypes
```

```
[5]: tripduration          int64
      starttime          datetime64[ns]
      stoptime           datetime64[ns]
      start station id     int64
      end station id       int64
      bikeid              int64
      birth year           int64
      gender              float64
      usertype_id          int64
      dtype: object
```

```
[6]: tripdata_df['month'] = tripdata_df['starttime'].dt.to_period(
      'M').dt.strftime("%Y-%m")
```

```
[7]: tripdata_df.head(2)
```

```
[7]:
```

	tripduration	starttime	stoptime	start station id \
0	211	2015-06-01 00:07:07	2015-06-01 00:10:39	88
1	834	2015-06-01 00:13:48	2015-06-01 00:27:43	5

	end station id	bikeid	birth year	gender	usertype_id	month
0	96	546	0	0.0	1	2015-06
1	12	487	1986	1.0	1	2015-06

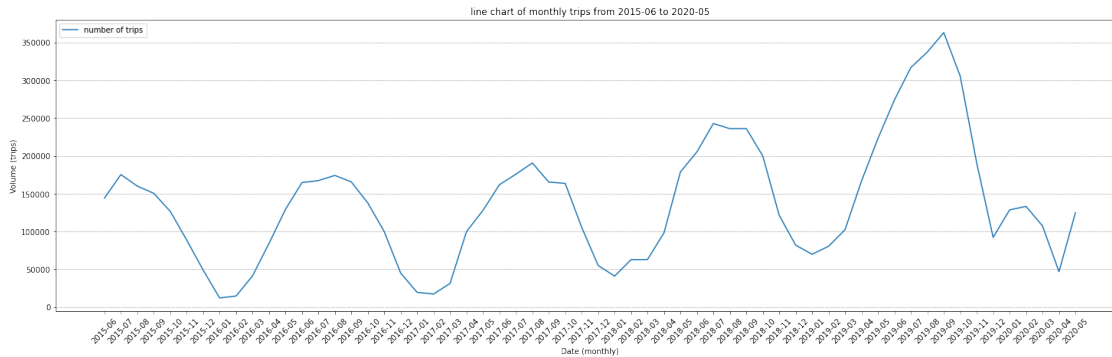
```
[8]: month_trip = tripdata_df['month'].value_counts().sort_index()
```

```
[74]: x = month_trip.index.tolist()
      y = month_trip.tolist()
```

```
[10]: # month_trip.plot()
      fig, ax = plt.subplots(figsize=(25, 7))
      ax.plot(x, y, label='number of trips')
      ax.set(xlabel='Date (monthly)', ylabel='Volume (trips)',
             title='line chart of monthly trips from 2015-06 to 2020-05')

      plt.grid(axis='y', color='grey', linestyle='--', linewidth=0.5)

      # You can specify a rotation for the tick labels in degrees or with keywords.
      plt.xticks(rotation='45')
      plt.legend(loc='upper left')
      plt.show()
```



2. Discuss any observed patterns, e.g., trend, seasonality, and shocks. Support your argument with necessary data, if possible. [2 pts]

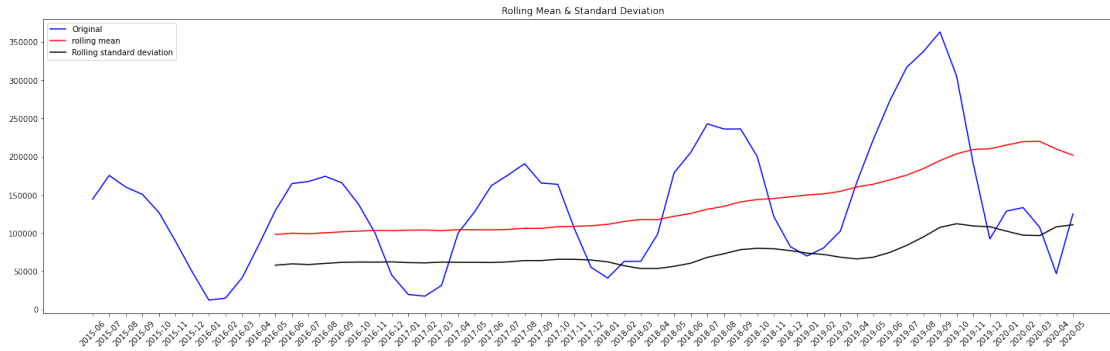
```
[11]: def plot_stationarity(timeseries):
    '''Plot Rolling Mean & Standard Deviation'''
    x = timeseries.index.tolist()

    # One year is used as a window, and the value of each time t is replaced by
    → the mean of the previous 12 months (including itself)
    # and the standard deviation is the same.
    rolmean = timeseries.rolling(window=12).mean()
    rolstd = timeseries.rolling(window=12).std()

    # plot rolling statistics:
    plt.subplots(figsize=(25, 7))
    plt.plot(x, timeseries.tolist(), color='blue', label='Original')
    plt.plot(x, rolmean.tolist(), color='red', label='rolling mean')
    plt.plot(x, rolstd.tolist(), color='black',
             label='Rolling standard deviation')

    plt.xticks(rotation='45')
    plt.legend(loc='upper left')
    plt.title('Rolling Mean & Standard Deviation')
    plt.show()
```

```
[12]: plot_stationarity(month_trip)
```



Through the above picture, we can clearly find that - The use of vehicles has a clear upward trend (until the beginning of 2020) - There is yearly seasonality (the peak is from July to September, and the trough is from November to next year February)

Check the Stationarity of time series data. Judging that the data is stable is often based on several statistics that are constant for time: - Constant mean - Constant variance - Time independent autocovariance

```
[13]: from statsmodels.tsa.stattools import adfuller
```

```
[14]: def test_stationarity(timeseries):
    '''Dickey-Fuller test'''

    print('Results of Dickey-Fuller Test:')
    dfctest = adfuller(timeseries, autolag='AIC')
    # The output value of #dfctest is Test Statistic, p-value, the number of
    ↪lags, the number of observations used
    # and the critical value under each confidence level.
    dfcoutput = pd.Series(dfctest[0:4],
                          index=['Test Statistic', 'p-value', '#Lags Used',
    ↪'Number of Observations Used'])
    for key, value in dfctest[4].items():
        dfcoutput['Critical value (%)' % key] = value

    print(dfcoutput)
```

```
[15]: test_stationarity(month_trip)
```

```
Results of Dickey-Fuller Test:
Test Statistic      0.758357
p-value             0.990933
#Lags Used          10.000000
Number of Observations Used 49.000000
Critical value (1%) -3.571472
Critical value (5%) -2.922629
```

```
Critical value (10%)          -2.599336
dtype: float64
```

The null hypothesis of the Augmented Dickey-Fuller is that there is a unit root, with the alternative that there is no unit root. If the pvalue is above a critical size, then we cannot reject that there is a unit root.(from [statsmodel official documentation](#))

- p-value
 - The p value is the evidence against a null hypothesis. The smaller the p-value, the stronger the evidence that you should reject the null hypothesis. In this case, the p-value is 0.99 which is a strong evidence that the null hypothesis can not be reject. In other words, the ADF test result shows that **there is a unit root**.
 - Nonstationarity can lead to spurious regression, which is an apparent relationship between variables that are not related in reality.
- Critical values for the test statistic at the 1 %, 5 %, and 10 % levels. Based on MacKinnon (2010).
 - The ADF test statistic value is greater than the critical value under each significant level, this result also means: the series is not stationary.

Note that: The ADF test does not prove nonstationarity; it fails to prove stationarity.

```
[16]: # log transform data to make data stationary on variance
month_trip_log = np.log10(month_trip)
# Difference data to make data stationary on mean (remove trend)
month_trip_stationary = month_trip_log.diff(periods=1)[1:]

[17]: test_stationarity(month_trip_stationary)
```

Results of Dickey-Fuller Test:

```
Test Statistic          -5.703544e+00
p-value                  7.579024e-07
#Lags Used               1.000000e+01
Number of Observations Used  4.800000e+01
Critical value (1%)      -3.574589e+00
Critical value (5%)      -2.923954e+00
Critical value (10%)     -2.600039e+00
dtype: float64
```

Difference log transform data to make data stationary on both mean and variance:

- p-value

In this case, the p-value is almost zero which is a strong evidence that the null hypothesis can be reject. In other words, the ADF test result shows that **there isn't a unit root**.
- Critical values for the test statistic

The ADF test statistic value is not greater than the critical value under each significant level, this result also means: the series is stationary.

3. Implement the following forecasting methods to forecast the monthly trips from 2019-06 to 2020-05. (Note that, when forecasting for month $t+1$, the history from 2015-06 up to month t are available.)

a. Provide your forecasts using a 3-month moving average. [1 pts]

```
[249]: from scipy.special import exp10
       from sklearn.metrics import mean_absolute_error as mae
```

```
[250]: # Define the mape metric
def mape(y_true, y_pred):
    y_true = np.array(y_true).astype(np.float64)
    y_pred = np.array(y_pred).astype(np.float64)
    if y_true.shape != y_pred.shape:
        raise ValueError(
            f"y_true and y_pred have different shape for "
            f"{y_true.shape} != {y_pred.shape}")
    return np.nanmean(np.abs((y_true - y_pred) / np.abs(y_true)))
```

```
[251]: def plot_forecast(df, other_info=''):
    fig, ax = plt.subplots(figsize=(25, 7))
    x = df.index.tolist()
    y_true = df['y_true'].tolist()
    y_pred = df['y_pred'].tolist()
    ax.plot(x, y_true, color='blue', label='Original')
    ax.plot(x, y_pred, color='red', label=other_info)

    ax.set_xlabel('Date (monthly)', fontsize=16)
    ax.set_ylabel('Volume (trips)', fontsize=16)
    ax.set_title(f'Forecasts using {other_info}', fontsize=20)

    ax.xaxis.set_tick_params(labelsize=16, rotation=45)
    ax.yaxis.set_tick_params(labelsize=16)
    ax.grid(axis='y', color='grey', linestyle='--', linewidth=0.5)

    # ax.xaxis.set_rotation('90')
    plt.legend(loc='upper left')
    plt.tight_layout()
    plt.show()
```

```
[252]: def moving_average_forecast(timeseries, sliding_window, start_forecast_date,
    ↪end_forecast_date):
    # CHECK THE DATA PREPARATION
    if month_trip.index[0] > str(pd.to_datetime(FORCAST_START_DATE) - pd.
    ↪offsets.MonthBegin(sliding_window)):
        raise Exception('In order to forecast month {}, the previous {} months_
    ↪data needed'.format(
        start_forecast_date, sliding_window))
```



```

# The value of each time t is replaced by the mean of the previous N_
↳ months(not including itself)
forecast_result = timeseries.rolling(window=3).mean().shift(1)
forecast_result = forecast_result.loc[forecast_result.index >=
↳ FORCAST_START_DATE]

result_df = pd.DataFrame({'y_true': timeseries, 'y_pred': forecast_result})

return result_df

```

```

[353]: def model_eval(y_ture, y_pred, start_date, end_date):
        y_ture = y_ture.loc[(y_ture.index >= start_date) &
                             (y_ture.index <= end_date)]
        y_pred = y_pred.loc[(y_pred.index >= start_date) &
                              (y_pred.index <= end_date)]
        print('MAPE: {}'.format(mape(y_ture, y_pred)))
        print('MAD: {}'.format(mae(y_ture, y_pred)))

```

```

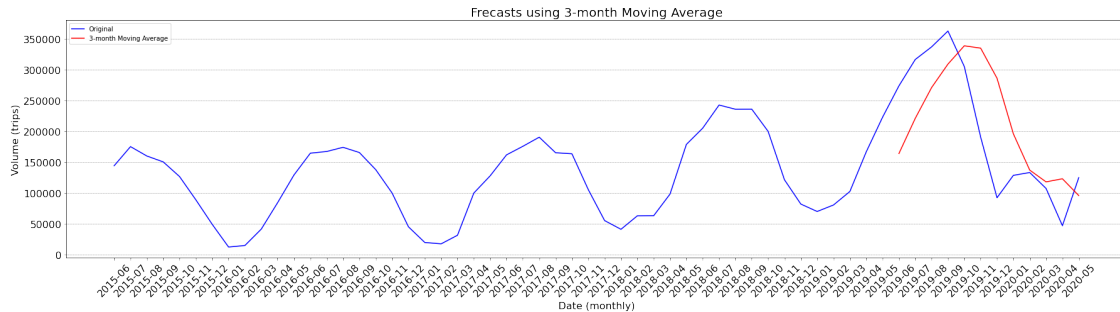
[320]: FORCAST_START_DATE = '2019-06'
        FORCAST_END_DATE = '2020-05'
        SLIDING_WINDOW = 3

```

```

[255]: result_df = moving_average_forecast(timeseries=month_trip,
                                           sliding_window=SLIDING_WINDOW,
                                           start_forecast_date=FORCAST_START_DATE,
                                           end_forecast_date=FORCAST_END_DATE)
        plot_forecast(result_df, other_info=f'{SLIDING_WINDOW}-month Moving Average')
        model_eval(result_df['y_true'],
↳ result_df['y_pred'], start_date=FORCAST_START_DATE, end_date=FORCAST_END_DATE)

```



MAPE: 0.5449525413832923

MAD: 73796.66666666667

b. Provide your forecasts using exponential smoothing with a smoothing constant $=0.5$. [1 pts]

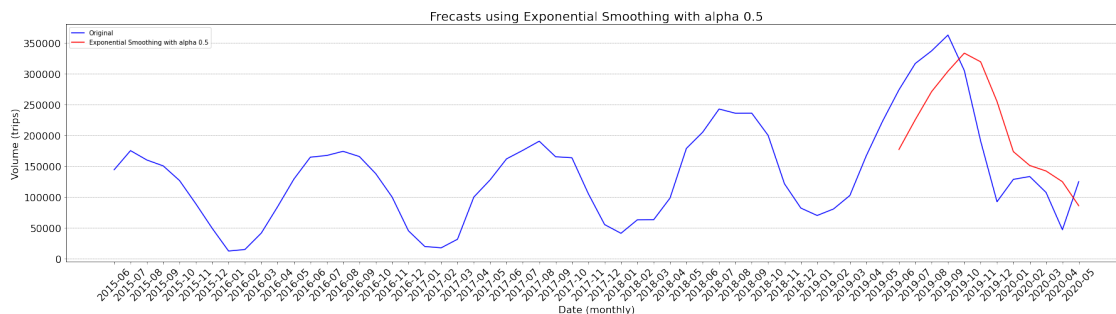
```
[256]: from statsmodels.tsa.api import SimpleExpSmoothing
```

```
[257]: def exponential_smoothing(timeseries, alpha, start_forecast_date,
    ↪end_forecast_date):
    """Provide forecasts using exponential smoothing with a smoothing
    ↪constant"""
    forecast_lst = pd.Series(pd.date_range(start_forecast_date,
    ↪end_forecast_date, freq='MS').strftime("%Y-%m"))
    forecast_value = []
    for forecast_month in forecast_lst:
        timeseries_train = timeseries.loc[timeseries.index < forecast_month].
    ↪tolist()
        SES_model = SimpleExpSmoothing(timeseries_train).
    ↪fit(smoothing_level=alpha, optimized=False)
        forecast_value.append(SES_model.forecast(1)[0])

    SES_result = pd.Series(forecast_value, index=forecast_lst)
    result_df = pd.DataFrame({'y_true': timeseries, 'y_pred': SES_result})
    return result_df
```

```
[258]: SES_ALPHA = 0.5
```

```
[259]: result_df = exponential_smoothing(timeseries=month_trip,
    alpha=SES_ALPHA,
    start_forecast_date=FORECAST_START_DATE,
    end_forecast_date=FORECAST_END_DATE)
plot_forecast(result_df, other_info=f'Exponential Smoothing with alpha
    ↪{SES_ALPHA}')
model_eval(result_df['y_true'],
    ↪result_df['y_pred'], start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)
```



MAPE: 0.5270872736829372

MAD: 70695.29033772778

c. Provide your forecasts using Holt's method with $\alpha=0.3$ and $\beta=0.1$. [1 pts]

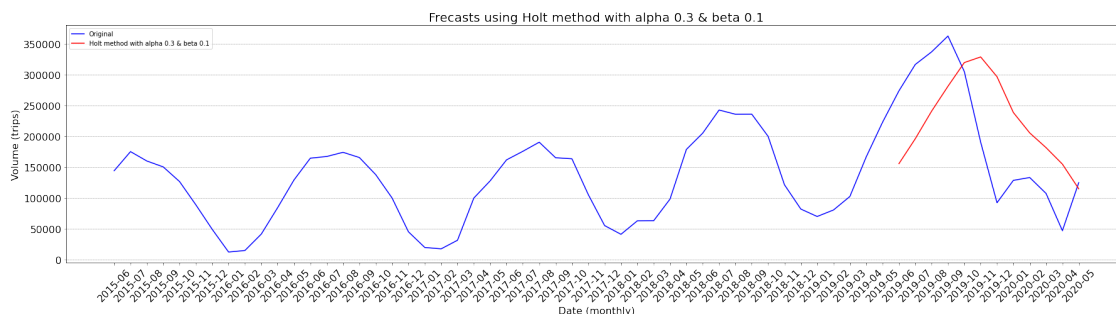
```
[260]: from statsmodels.tsa.api import Holt
```

```
[261]: def holt_method(timeseries, alpha, beta, start_forecast_date,
    ↪end_forecast_date):
    """Provide forecasts using Holt's method with alpha and beta"""
    forecast_lst = pd.Series(pd.date_range(
        start_forecast_date, end_forecast_date, freq='MS').strftime("%Y-%m"))
    forecast_value = []
    for forecast_month in forecast_lst:
        timeseries_train = timeseries.loc[timeseries.index < forecast_month].
    ↪tolist()
        Holter_model = Holt(timeseries_train).fit(
            smoothing_level=alpha, smoothing_slope=beta, optimized=False)
        forecast_value.append(Holter_model.forecast(1)[0])

    Holter_result = pd.Series(forecast_value, index=forecast_lst)
    result_df = pd.DataFrame({'y_true': timeseries, 'y_pred': Holter_result})
    return result_df
```

```
[262]: HOLT_ALPHA = 0.3
    HOLT_BETA = 0.1
```

```
[263]: result_df = holt_method(timeseries=month_trip,
    alpha=HOLT_ALPHA,
    beta=HOLT_BETA,
    start_forecast_date=FORCAST_START_DATE,
    end_forecast_date=FORCAST_END_DATE)
plot_forecast(result_df, other_info=f'Holt method with alpha {HOLT_ALPHA} & beta_
    ↪{HOLT_BETA}')
model_eval(result_df['y_true'],
    ↪result_df['y_pred'], start_date=FORCAST_START_DATE, end_date=FORCAST_END_DATE)
```



MAPE: 0.7338497610269354
MAD: 95866.41657870775

d. Plot the forecasts from all three methods above in the line chart, together with the actual trips for the periods 2019-06 to 2020-05. [1 pts]

```
[269]: def plot_all_method(full_results):
    fig, ax = plt.subplots(figsize=(25, 7))
    x = full_results.index.tolist()

    ax.plot(x, full_results['y_true'].tolist(), color='blue', label='Original')
    ax.plot(x, full_results['y_pred_ma'].tolist(),
    ↪color='red', label='y_pred_ma')
    ax.plot(x, full_results['y_pred_es'].tolist(),
    ↪color='orange', label='y_pred_es')
    ax.plot(x, full_results['y_pred_hm'].tolist(),
    ↪color='green', label='y_pred_hm')

    ax.set_xlabel('Date (monthly)', fontsize=16)
    ax.set_ylabel('Volume (trips)', fontsize=16)
    ax.set_title(f'Forecasts using all three methods', fontsize=20)

    ax.xaxis.set_tick_params(labelsize=16, rotation=45)
    ax.yaxis.set_tick_params(labelsize=16)
    ax.grid(axis='y', color='grey', linestyle='--', linewidth=0.5)

    # ax.xaxis.set_rotation('45')
    plt.legend(loc='upper left')
    plt.tight_layout()
    plt.show()
```

```
[271]: ma_result = moving_average_forecast(
    timeseries=month_trip, sliding_window=SLIDING_WINDOW,
    start_forecast_date=FORECAST_START_DATE,
    end_forecast_date=FORECAST_END_DATE)
ma_result.rename(columns={'y_pred': 'y_pred_ma'}, inplace=True)

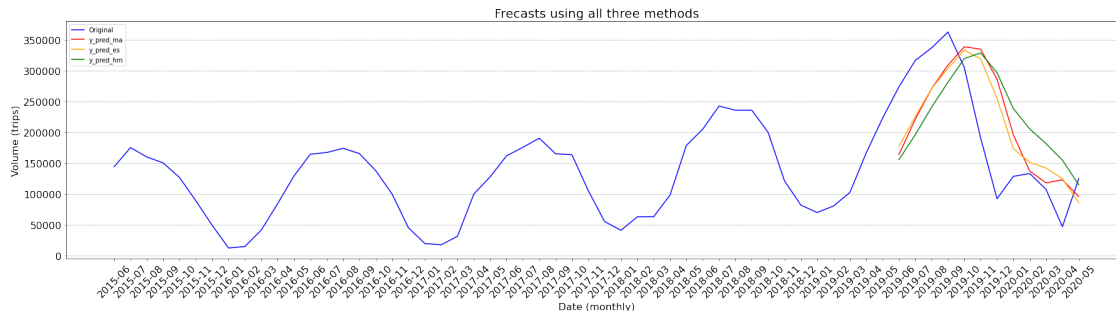
es_result = exponential_smoothing(
    timeseries=month_trip, alpha=SES_ALPHA,
    start_forecast_date=FORECAST_START_DATE,
    end_forecast_date=FORECAST_END_DATE)
es_result.rename(columns={'y_pred': 'y_pred_es'}, inplace=True)

hm_result = holt_method(
    timeseries=month_trip, alpha=HOLT_ALPHA, beta=HOLT_BETA,
    start_forecast_date=FORECAST_START_DATE,
    end_forecast_date=FORECAST_END_DATE)
hm_result.rename(columns={'y_pred': 'y_pred_hm'}, inplace=True)

full_results = pd.merge(ma_result, es_result['y_pred_es'],
    left_index=True, right_index=True)
```

```
full_results = pd.merge(full_results, hm_result['y_pred_hm'],
                        left_index=True, right_index=True)

plot_all_method(full_results)
```



e. Evaluate the above forecasting methods using MAD (Mean Absolute Deviation) and MAPE, respectively. [1 pts]

```
[266]: print('***** MOVING AVERAGE *****')
model_eval(full_results['y_true'], full_results['y_pred_ma'],
            start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)

print('\n***** EXPONENTIAL SMOOTHING *****')
model_eval(full_results['y_true'], full_results['y_pred_es'],
            start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)

print('\n***** HOLT METHOD *****')
model_eval(full_results['y_true'], full_results['y_pred_hm'],
            start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)
```

```
***** MOVING AVERAGE *****
MAPE: 0.5449525413832923
MAD: 73796.66666666667
```

```
***** EXPONENTIAL SMOOTHING *****
MAPE: 0.5270872736829372
MAD: 70695.29033772778
```

```
***** HOLT METHOD *****
MAPE: 0.7338497610269354
MAD: 95866.41657870775
```

f. Discuss the performance of the forecasting methods, e.g., any suggestions for improvement. [1 pts]

- We can observe that the three methods have the same delay relative to the real situation.

This delay is most likely due to seasonality.

- Suggestions for improvement: transform the original data, using the difference and log method. Carry out stationarity test to verify the stationarity of the time series.
- In this case(using stationary data), we need to convert the forecast on differenced data to non-differenced data forecast.

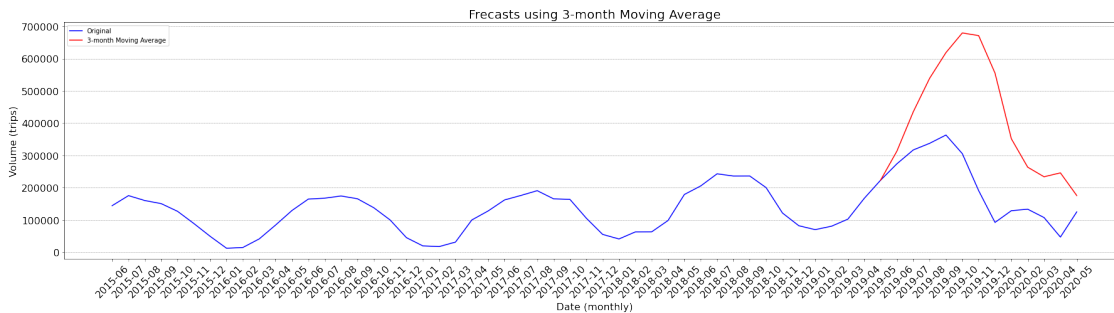
```
[221]: def recover_origin(ts, ts_log, ts_stationary_pred, forecast_start_date,
    ↪pred_type):
    """Predicted values need to be restored by the relevant inverse_
    ↪transformation."""
    ts_stationary_pred = ts_stationary_pred.dropna()
    if pred_type == "daily":
        origin_index = (pd.to_datetime(forecast_start_date) - pd.offsets.
    ↪DateOffset(1)).strftime("%Y-%m-%d")
        origin_num = ts_log[origin_index]
    elif pred_type == 'monthly':
        origin_index = (pd.to_datetime(forecast_start_date) - pd.offsets.
    ↪MonthBegin(1)).strftime("%Y-%m")
        origin_num = ts_log[origin_index]
    else:
        raise Exception
    # Recover the first-order difference
    diff_restored = pd.Series(origin_num, index=[origin_index]).
    ↪append(ts_stationary_pred).cumsum()

    # recover the log transformation
    log_recover = 10 ** diff_restored
    log_recover = log_recover.astype(int)

    # concatenate the result
    result = pd.concat([ts, log_recover], axis=1)
    result.columns = ['y_true', 'y_pred']
    return result
```

```
[163]: station_result_df = moving_average_forecast(timeseries=month_trip_stationary,
    sliding_window=SLIDING_WINDOW,
    start_forecast_date=FORECAST_START_DATE,
    end_forecast_date=FORECAST_END_DATE)

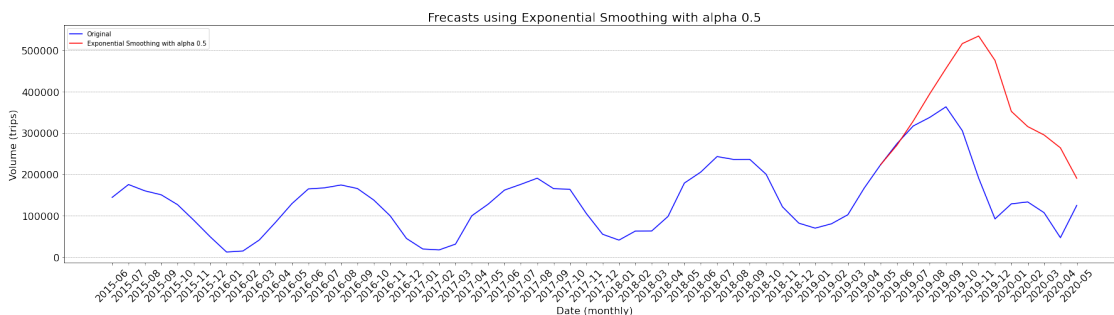
result_df = recover_origin(month_trip, month_trip_log,
    ↪station_result_df['y_pred'], FORECAST_START_DATE, pred_type='monthly')
plot_forecast(result_df, other_info=f'{SLIDING_WINDOW}-month Moving Average')
model_eval(result_df['y_true'],
    ↪result_df['y_pred'], start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)
```



MAPE: 1.5944921516090218

MAD: 221843.16666666666

```
[267]: station_result_df = exponential_smoothing(timeseries=month_trip_stationary,
                                                alpha=ALPHA,
                                                start_forecast_date=FORCAST_START_DATE,
                                                end_forecast_date=FORCAST_END_DATE)
result_df = recover_origin(month_trip, month_trip_log,
                           station_result_df['y_pred'],
                           FORCAST_START_DATE, pred_type='monthly')
plot_forecast(result_df,
              other_info=f'Exponential Smoothing with alpha {SES_ALPHA}')
model_eval(result_df['y_true'],
           result_df['y_pred'], start_date=FORCAST_START_DATE, end_date=FORCAST_END_DATE)
```



MAPE: 1.4285772827330065

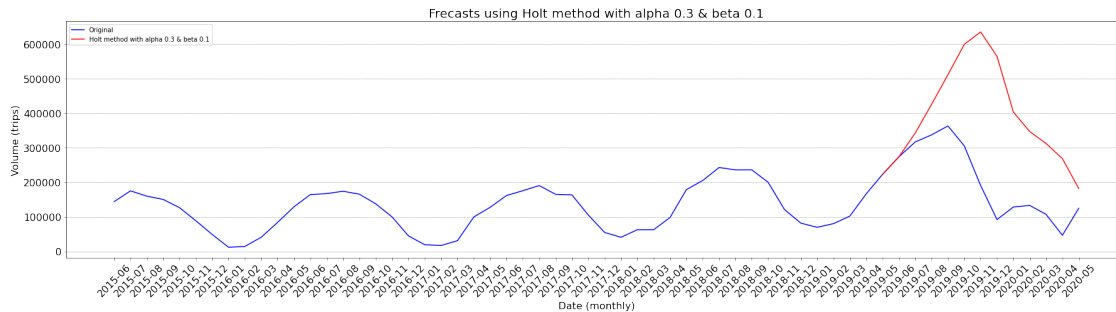
MAD: 164884.41666666666

```
[268]: station_result_df = holt_method(timeseries=month_trip_stationary,
                                       alpha=ALPHA,
                                       beta=BETA,
                                       start_forecast_date=FORCAST_START_DATE,
                                       end_forecast_date=FORCAST_END_DATE)
```

```

result_df = recover_origin(month_trip, month_trip_log,
    ↳station_result_df['y_pred'], FORCAST_START_DATE, pred_type='monthly')
plot_forecast(result_df, other_info=f'Holt method with alpha {HOLT_ALPHA} & beta {HOLT_BETA}')
model_eval(result_df['y_true'],
    ↳result_df['y_pred'], start_date=FORCAST_START_DATE, end_date=FORCAST_END_DATE)

```



MAPE: 1.6684936030127824

MAD: 203917.75

```

[272]: ma_station_result = moving_average_forecast(
    timeseries=month_trip_stationary, sliding_window=SLIDING_WINDOW,
    start_forecast_date=FORCAST_START_DATE,
    end_forecast_date=FORCAST_END_DATE)
ma_result = recover_origin(month_trip, month_trip_log,
    ↳ma_station_result['y_pred'], FORCAST_START_DATE, pred_type='monthly')
ma_result.rename(columns={'y_pred': 'y_pred_ma'}, inplace=True)

es_station_result = exponential_smoothing(
    timeseries=month_trip_stationary, alpha=SES_ALPHA,
    start_forecast_date=FORCAST_START_DATE,
    end_forecast_date=FORCAST_END_DATE)
es_result = recover_origin(month_trip, month_trip_log,
    ↳es_station_result['y_pred'], FORCAST_START_DATE, pred_type='monthly')
es_result.rename(columns={'y_pred': 'y_pred_es'}, inplace=True)

hm_station_result = holt_method(
    timeseries=month_trip_stationary, alpha=HOLT_ALPHA, beta=HOLT_BETA,
    start_forecast_date=FORCAST_START_DATE,
    end_forecast_date=FORCAST_END_DATE)
hm_result = recover_origin(month_trip, month_trip_log,
    ↳hm_station_result['y_pred'], FORCAST_START_DATE, pred_type='monthly')
hm_result.rename(columns={'y_pred': 'y_pred_hm'}, inplace=True)

full_results = pd.merge(ma_result, es_result['y_pred_es'],

```

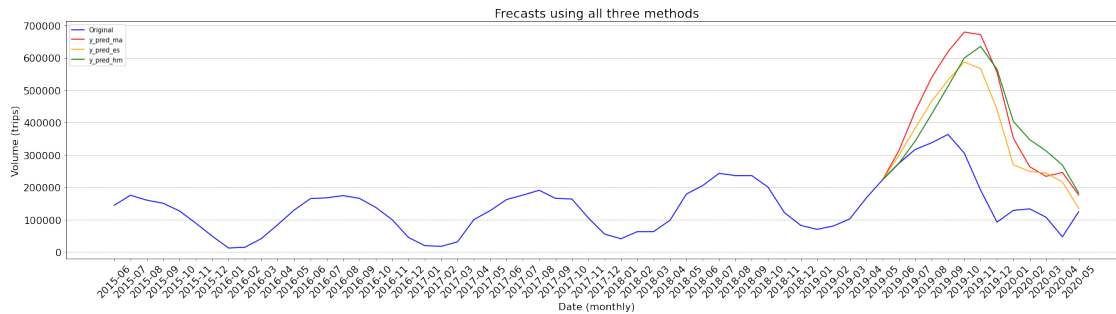


```

        left_index=True, right_index=True)
full_results = pd.merge(full_results, hm_result['y_pred_hm'],
        left_index=True, right_index=True)

plot_all_method(full_results)

```



```

[273]: print('***** MOVING AVERAGE *****')
model_eval(full_results['y_true'], full_results['y_pred_ma'],
        start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)

print('\n***** EXPONENTIAL SMOOTHING *****')
model_eval(full_results['y_true'], full_results['y_pred_es'],
        start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)

print('\n***** HOLT METHOD *****')
model_eval(full_results['y_true'], full_results['y_pred_hm'],
        start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)

```

***** MOVING AVERAGE *****

MAPE: 1.5944921516090218

MAD: 221843.16666666666

***** EXPONENTIAL SMOOTHING *****

MAPE: 1.2278635356494478

MAD: 163576.33333333334

***** HOLT METHOD *****

MAPE: 1.6684936030127824

MAD: 203917.75

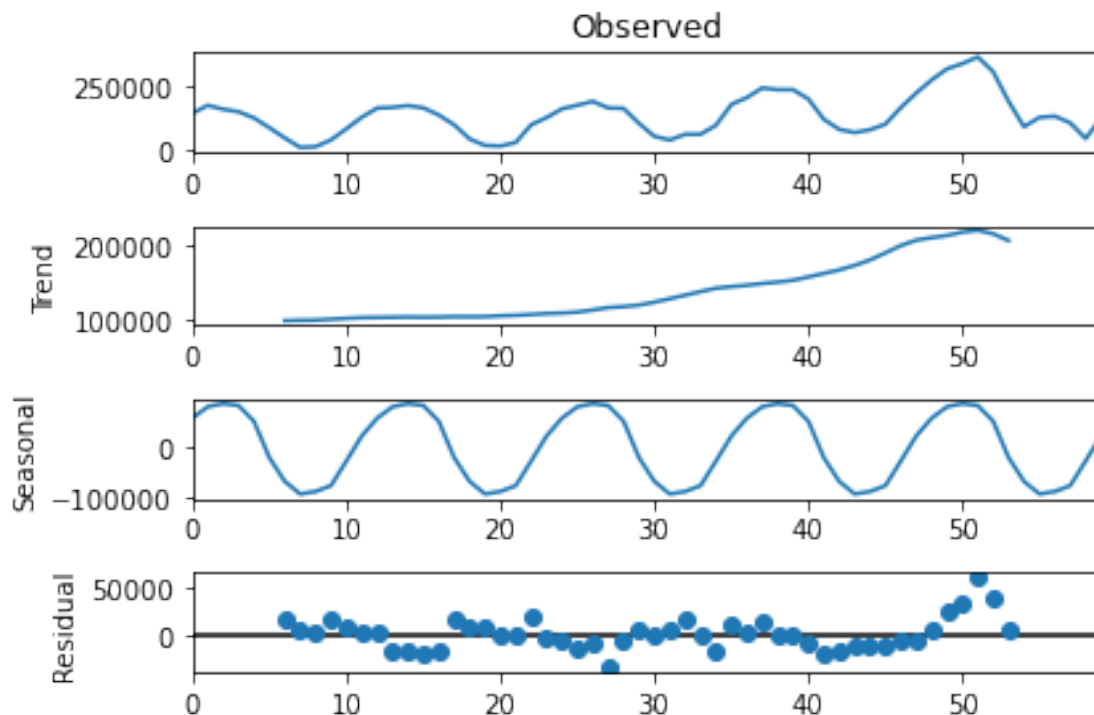
- Through the analysis of the results, the result of using the stationary sequence to predict is worse than that of the original sequence.
- But it does eliminate the delay that appeared before.
- Part of the reason why the result became worse is that none of the three methods can completely eliminate the upward trend of the sequence, and the use of exponents when restoring the original sequence may amplify this defect.

4. Use the histories from 2015-06 to 2019-05 to forecast the monthly trips from 2019-06 to 2020-05 using Holt-Winters' method and ARIMA.

a. Show the time series decomposition for periods 2015-06 to 2019-05 to identify trend and seasonality, if any. [1 pt]

```
[38]: from statsmodels.tsa.seasonal import seasonal_decompose
```

```
[39]: result = seasonal_decompose(month_trip.tolist(), model='additive', period=12)
result.plot()
plt.show()
```



The results here are consistent with the previous analysis: - The use of vehicles has a clear upward trend (until the beginning of 2020) - There is yearly seasonality (the peak is from July to September, and the trough is from November to next year February)

b. Provide your forecasts using Holt-Winters' method. [2 pts]

```
[40]: from statsmodels.tsa.api import ExponentialSmoothing
```

```
[321]: def holt_winters_method(timeseries, start_forecast_date, end_forecast_date):
timeseries_train = timeseries.loc[timeseries.index < start_forecast_date]
timeseries_forecast = timeseries.loc[timeseries.index >=
start_forecast_date]
step_num = len(timeseries_forecast)
```

```

# to silence the ValueError: No frequency information was provided, so
↳inferred frequency MS will be used.
timeseries_train.index = pd.DatetimeIndex(
    timeseries_train.index.values, freq='MS')

ES_model = ExponentialSmoothing(
    timeseries_train, trend='add', seasonal='add', seasonal_periods=12).
↳fit()
print(ES_model.summary())

ES_model = ES_model.forecast(step_num)
ES_model.index = ES_model.index.strftime("%Y-%m")

result_df = pd.DataFrame({'y_true': timeseries, 'y_pred': ES_model})

return result_df

```

```

[322]: result_df = holt_winters_method(timeseries=month_trip,
                                     start_forecast_date=FORECAST_START_DATE,
                                     end_forecast_date=FORECAST_END_DATE)
plot_forecast(result_df, other_info="Holt-Winters' Method")
model_eval(result_df['y_true'], result_df['y_pred'],
            start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)

```

ExponentialSmoothing Model Results

```

=====
Dep. Variable:                endog    No. Observations:                48
Model:                        ExponentialSmoothing    SSE                6994258936.595
Optimized:                    True      AIC                934.263
Trend:                        Additive    BIC                964.203
Seasonal:                    Additive    AICC                957.850
Seasonal Periods:            12      Date:                Tue, 15 Sep 2020
Box-Cox:                     False      Time:                01:15:01
Box-Cox Coeff.:              None
=====

```

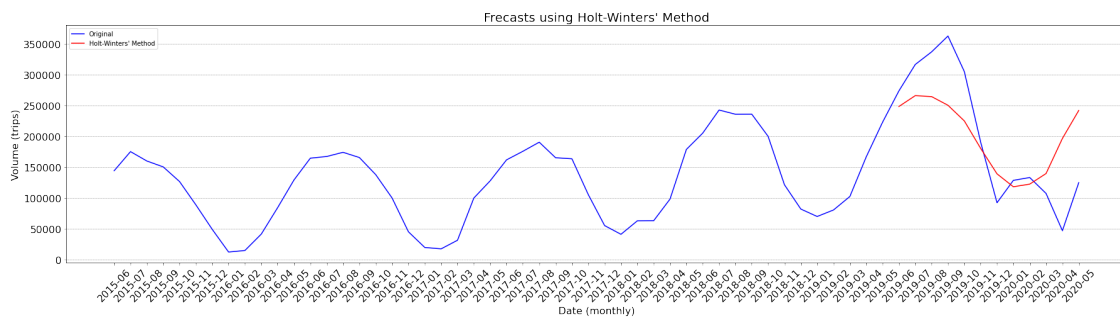
```

=
                                coeff                code                optimized
-----
-
smoothing_level                0.5525095                alpha
True
smoothing_slope                2.3737e-08                beta
True
smoothing_seasonal            0.000000                gamma
True
initial_level                  1.5116e+05                1.0

```

True		
initial_slope	2181.9353	b.0
True		
initial_seasons.0	-11469.657	s.0
True		
initial_seasons.1	3958.7294	s.1
True		
initial_seasons.2	220.38007	s.2
True		
initial_seasons.3	-16040.221	s.3
True		
initial_seasons.4	-43629.340	s.4
True		
initial_seasons.5	-90502.400	s.5
True		
initial_seasons.6	-1.3405e+05	s.6
True		
initial_seasons.7	-1.5724e+05	s.7
True		
initial_seasons.8	-1.5514e+05	s.8
True		
initial_seasons.9	-1.4012e+05	s.9
True		
initial_seasons.10	-85277.338	s.10
True		
initial_seasons.11	-41937.125	s.11
True		

-



MAPE: 0.5174405163453567

MAD: 59948.10268166926

c. Identification of best fit ARIMA model. Explain the resulting model, e.g., any (seasonal) differencing. [2 pts]

```
[323]: import pmdarima as pm
```

```
[324]: model = pm.auto_arima(month_trip, seasonal = True, suppress_warnings=True)
print(model.summary())
```

```

                                SARIMAX Results
=====
Dep. Variable:                  y      No. Observations:                  60
Model:                        SARIMAX(2, 0, 0)  Log Likelihood                -705.009
Date:                        Tue, 15 Sep 2020    AIC                        1418.019
Time:                        01:15:21          BIC                        1426.396
Sample:                        0              HQIC                       1421.296
                                - 60
Covariance Type:              opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
intercept    3.113e+04    9008.530      3.455      0.001     1.35e+04     4.88e+04
ar.L1         1.4408         0.080     18.018      0.000         1.284         1.598
ar.L2        -0.6638         0.091     -7.285      0.000        -0.842        -0.485
sigma2       9.224e+08         0.145    6.35e+09      0.000     9.22e+08     9.22e+08
=====
===
Ljung-Box (Q):                  46.69    Jarque-Bera (JB):
13.92
Prob(Q):                        0.22    Prob(JB):
0.00
Heteroskedasticity (H):         5.36    Skew:
0.97
Prob(H) (two-sided):            0.00    Kurtosis:
4.36
=====
===

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

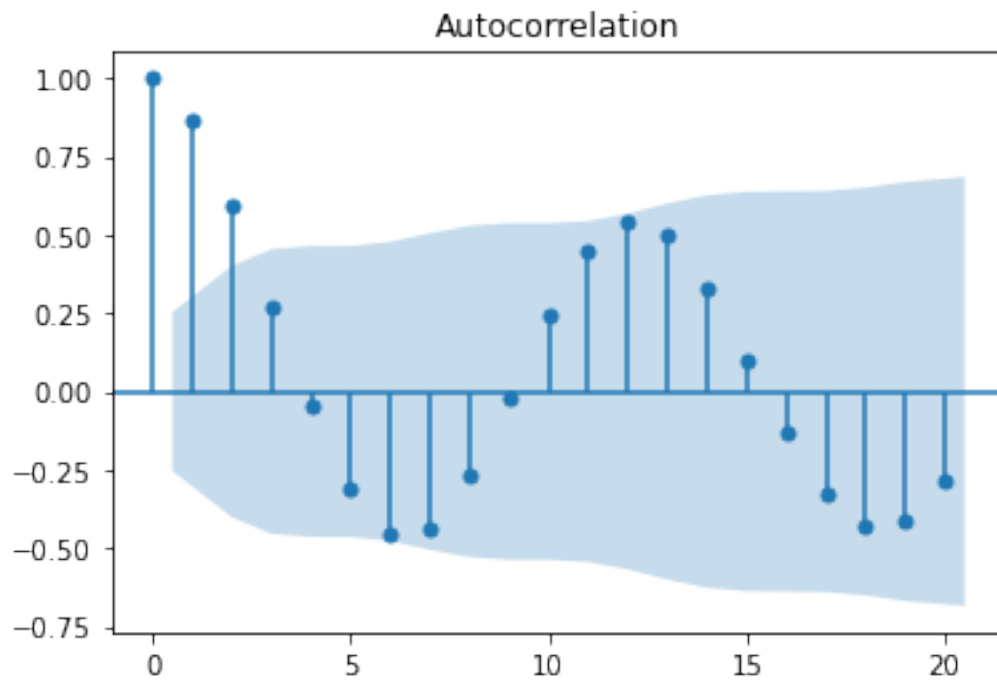
[2] Covariance matrix is singular or near-singular, with condition number 3.09e+25. Standard errors may be unstable.

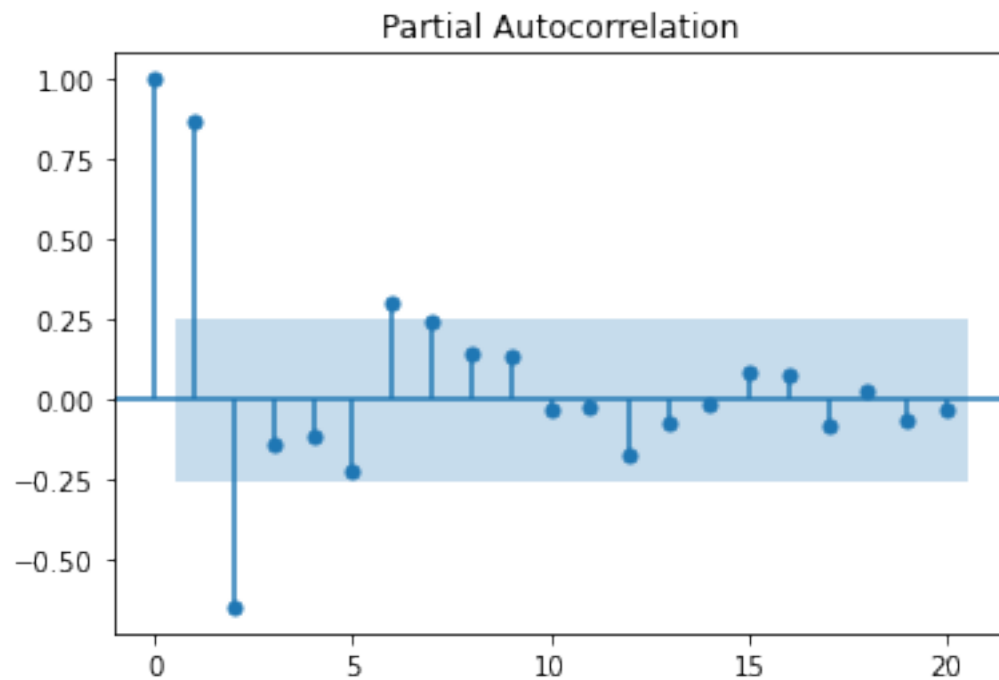
- A model with (only) two AR terms would be specified as an ARIMA of order (2,0,0).
- The SARIMA extension of ARIMA that explicitly models the seasonal element in univariate data. And there is no seasonal order taken into the model.

d. Plot ACF and PACF of fitted residuals to verify whether there is MA/AR effect left. [1 pt]

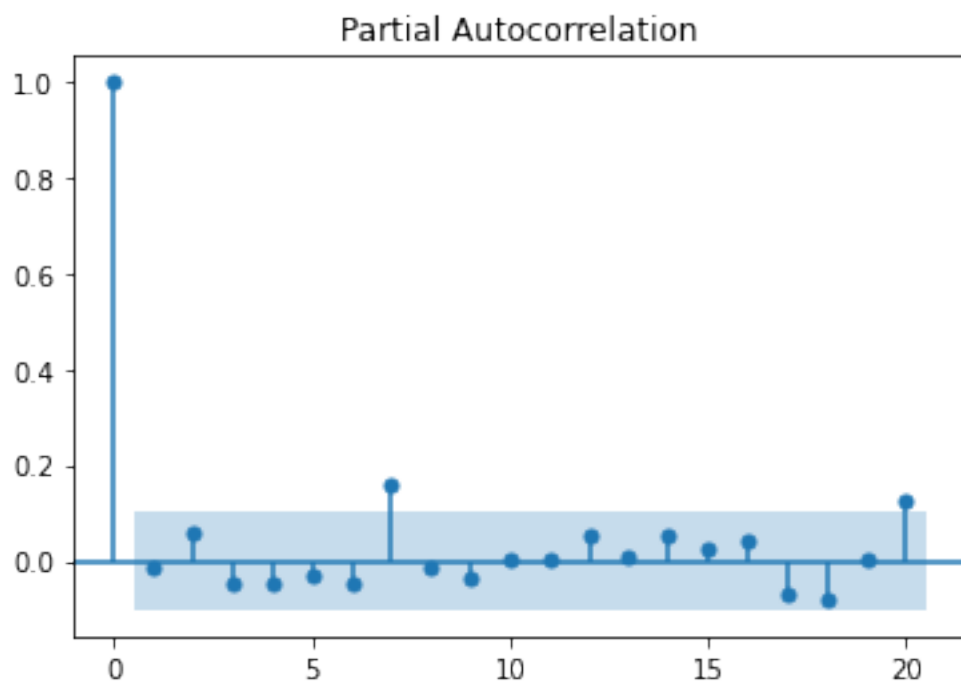
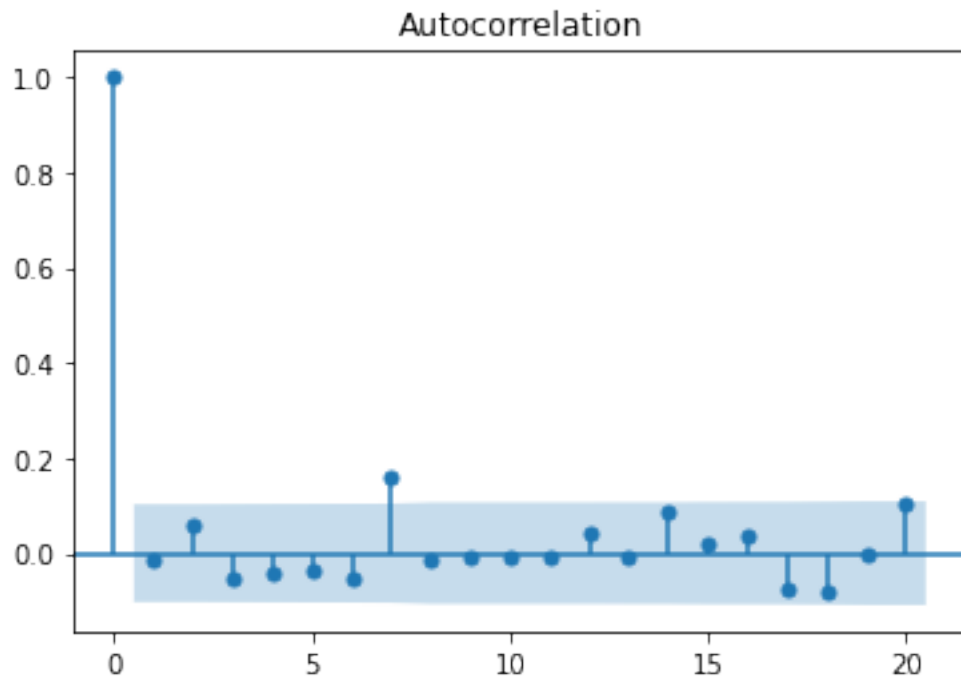
```
[45]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
[276]: # plot original time series
plot_acf(month_trip.tolist(), lags=20)
plt.show()
# set method to silence RuntimeWarning
plot_pacf(month_trip.tolist(), lags=20, method='ywm')
plt.show()
```





```
[274]: # plot residual
plot_acf(model.resid(), lags=20)
plt.show()
plot_pacf(model.resid(), lags=20, method='ywmm')
plt.show()
```



e. Forecast the trips for 2019-06 to 2020-05 using the best fit ARIMA model and plot the predictions with 95% confidence intervals. [2 pts]


```
[47]: def arima_forecast(timeseries, start_forecast_date, end_forecast_date):

    # Load/split data
    timeseries_train = timeseries.loc[timeseries.index < start_forecast_date]
    timeseries_forecast = timeseries.loc[(timeseries.index >=
    ↪start_forecast_date) &
                                (timeseries.index <=
    ↪end_forecast_date)]
    step_num = len(timeseries_forecast)

    # Fit model
    model = pm.auto_arima(timeseries_train, seasonal=True,
                          m=12, suppress_warnings=True) # m=seasonal length

    # make forecasts
    # predict N steps into the future
    pred, conf = model.predict(step_num, return_conf_int=True, alpha=0.05)

    model_df = pd.DataFrame()
    model_df['lower_bounds'] = [i[0] for i in conf]
    model_df['upper_bounds'] = [i[1] for i in conf]
    model_df['y_pred'] = list(pred)
    model_df.index = pd.Series(pd.date_range(
        start_forecast_date, end_forecast_date, freq='MS').strftime("%Y-%m"))

    result_df = pd.merge(left=pd.DataFrame(month_trip).rename(columns={'month':
    ↪'y_true'}),
                          right=model_df, how='left', right_index=True,
    ↪left_index=True)

    return result_df, model
```

```
[280]: def arima_plot(arima_result, title):
    fig, ax = plt.subplots(figsize=(25, 7))
    x = arima_result.index.tolist()
    ax.plot(x, arima_result['y_true'].tolist(), color='blue',
            label='Original')
    ax.plot(x, arima_result['y_pred'].tolist(), color='red',
            label='prediction')
    ax.plot(x, arima_result['lower_bounds'].tolist(), color='orange',
            label='Lower Bounds')
    ax.plot(x, arima_result['upper_bounds'].tolist(), color='green',
            label='Upper Bounds')

    ax.
    ↪fill_between(x, arima_result['lower_bounds'], arima_result['upper_bounds'], facecolor='silver')
```

```

ax.set_xlabel('Date (monthly)', fontsize=16)
ax.set_ylabel('Volume (trips)', fontsize=16)
ax.set_title(title, fontsize=20)

ax.xaxis.set_tick_params(labelsize=16, rotation=45)
ax.yaxis.set_tick_params(labelsize=16)
ax.grid(axis='y', color='grey', linestyle='--', linewidth=0.5)

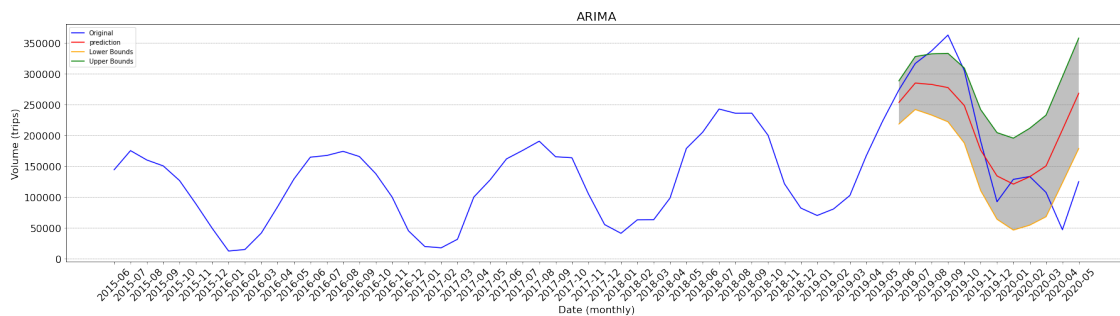
# ax.xaxis.set_rotation('45')
plt.legend(loc='upper left')
plt.tight_layout()
plt.show()

```

```

[281]: result_df, model = arima_forecast(timeseries=month_trip,
                                         start_forecast_date=FORECAST_START_DATE,
                                         end_forecast_date=FORECAST_END_DATE)
arima_plot(result_df, title='ARIMA')
model_eval(result_df['y_true'], result_df['y_pred'],
           start_date=FORECAST_START_DATE, end_date=FORECAST_END_DATE)

```



MAPE: 0.5310480887395453

MAD : 55215.56640086445

5. Focus on the trip data from 2019-01 to 2019-12 only. Now the client of MSBA & Company wants to understand the key factors that affects the total daily ridership in the studied periods. Please explore auxiliary data sets and discuss your findings. [3 pts] [Hint: Weather data: https://www.meteoblue.com/en/weather/archive/era5/boston_united-states-of-america_4930956 . Students are also encouraged to explore other auxiliary data source]

Features: - Weather Dataset - Temperature: basic indicator - Precipitation: an indicator of whether rain or not - Wind speed: an indicator of whether suitable for taking scooter - Date-related Dataset - holiday - etc.

```
[332]: # Set hyper parameter
TRAIN_START_DATE = '2019-01-01'
FORCAST_START_DATE = '2019-10-01'
FORCAST_END_DATE = '2019-12-31'

[333]: tripdata_df['datetime'] = tripdata_df['starttime'].dt.strftime("%Y-%m-%d")
tripdata_used = tripdata_df.loc[(tripdata_df['datetime'] >= TRAIN_START_DATE) &
                                (tripdata_df['datetime'] <= FORCAST_END_DATE)]

[334]: daily_trip = tripdata_used['datetime'].value_counts().sort_index()

[335]: daily_trip_df = pd.DataFrame(daily_trip).rename(columns={'datetime': 'y_true'})
daily_trip_df['datetime'] = pd.to_datetime(daily_trip_df.index)
daily_trip_df = daily_trip_df.reset_index(drop=True)

[336]: weather_data = read_from_csv('weather_data.csv', PROCESSED_DATA_DIR,
                                   parse_dates=['datetime'])
weather_data['datetime'] = weather_data['datetime']

[337]: dataset = pd.merge(left=daily_trip_df, right=weather_data,
                          on='datetime', how='inner')

[338]: # Add Date-related Features
dataset['day'] = dataset['datetime'].dt.strftime("%d").astype(int)
dataset['dayofweek'] = dataset['datetime'].dt.dayofweek
dataset['is_workday'] = dataset['dayofweek'].apply(
    lambda x: 1 if x <= 4 else 0)

[339]: import holidays
from datetime import date

[340]: us_holidays = holidays.UnitedStates()
dataset['is_holiday'] = dataset['datetime'].apply(
    lambda x: 1 if x in us_holidays else 0)

[341]: dataset.head()
```

	y_true	datetime	Temperature_Minimum	Temperature_Maximum	\
0	1305	2019-01-01	2.348312	13.128311	
1	2632	2019-01-02	-2.821688	1.618311	
2	3005	2019-01-03	-1.311689	6.968312	
3	3397	2019-01-04	-0.161688	7.998312	
4	786	2019-01-05	1.728312	5.048312	

	Temperature_Mean	RelativeHumidity_Minimum	RelativeHumidity_Maximum	\
0	7.768727	49.318592	95.652000	
1	-0.883772	31.418629	67.223060	

2	2.654561	57.154278	88.330734
3	3.625811	48.568703	73.006860
4	3.663312	70.087240	93.792020

	RelativeHumidity_Mean	MeanSeaLevelPressure_Minimum \
0	71.862465	997.7
1	48.663450	1018.7
2	71.750340	1009.3
3	63.310513	1009.5
4	88.568260	998.5

	MeanSeaLevelPressure_Maximum ...	WindSpeed_Maximum	WindSpeed_Mean \
0	1018.1 ...	31.780067	23.699722
1	1027.7 ...	14.512064	10.105483
2	1022.3 ...	16.793140	11.653485
3	1015.1 ...	18.075441	12.926744
4	1008.7 ...	17.727943	10.884510

	WindDirection Dominant_None	GeopotentialHeight_Minimum \
0	261.37690	-18
1	339.22775	149
2	245.04233	75
3	225.24014	77
4	352.37183	-12

	GeopotentialHeight_Maximum	GeopotentialHeight_Mean	day	dayofweek \
0	144	55.916668	1	1
1	216	193.750000	2	2
2	176	109.000000	3	3
3	120	98.916664	4	4
4	70	24.541666	5	5

	is_workday	is_holiday
0	1	1
1	1	0
2	1	0
3	1	0
4	0	0

[5 rows x 33 columns]

SARIMAX is like a complicated version of ARIMA. Optimal values of p, d, and q can be searched via a loop and grid search, as well as the seasonal values for p, d, and q. And also there could add more parameters in the SARIMAX function.

```
[342]: import statsmodels.api as sm
```

```
[371]: def sarimax_forecast(df, exog_lst, start_forecast_date, end_forecast_date):

    # Load/split data
    df_train = df.loc[df['datetime'] < start_forecast_date]
    df_forecast = df.loc[(df['datetime'] >= start_forecast_date) &
                          (df['datetime'] <= end_forecast_date)]

    step_num = len(df_forecast)

    # Split the observed time-series and exogenous regressors
    endog_train = df_train['y_true']
    endog_forecast = df_forecast['y_true']
    exog_train = df_train[exog_lst]
    exog_forecast = df_forecast[exog_lst]

    # Fit model
    # Make forecasts
    if len(exog_lst):
        model = sm.tsa.statespace.SARIMAX(endog=endog_train, exog=exog_train,
                                          enforce_stationarity=False,
                                          enforce_invertibility=False)

        model = model.fit(maxiter=1000, warn_convergence=False)
        forecast = model.get_forecast(steps=step_num, exog=exog_forecast)
    else:
        model = sm.tsa.statespace.SARIMAX(endog=endog_train,
                                          enforce_stationarity=False,
                                          enforce_invertibility=False)

        model = model.fit(maxiter=1000, warn_convergence=False)
        forecast = model.get_forecast(steps=step_num)

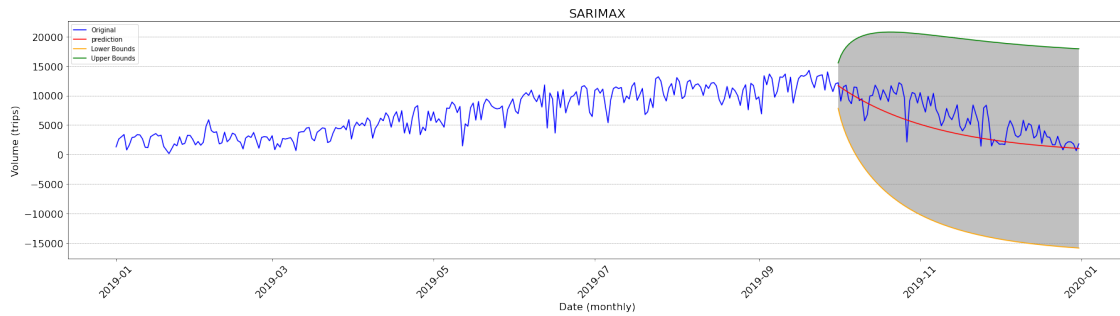
    forecast_interval = forecast.conf_int()
    y_pred = forecast.predicted_mean
    pred_all = pd.concat([y_pred, forecast_interval], axis=1)
    pred_all.columns = ['y_pred', 'lower_bounds', 'upper_bounds']
    pred_all['datetime'] = pd.date_range(
        start_forecast_date, end_forecast_date, freq='D')

    result_df = pd.merge(left=df[['y_true', 'datetime']],
                        right=pred_all, how='left', on='datetime')
    result_df.index = result_df['datetime']

    return result_df, model

[372]: result_df, model = sarimax_forecast(df=dataset, exog_lst=[],
                                          start_forecast_date=FORECAST_START_DATE,
                                          end_forecast_date=FORECAST_END_DATE)
arima_plot(result_df, title='SARIMAX')
```

```
model_eval(result_df['y_true'], result_df['y_pred'],
           start_date=FORCAST_START_DATE, end_date=FORCAST_END_DATE)
```



MAPE: 0.40792774715829755

MAD: 2389.541413092803

```
[355]: print(model.summary())
```

SARIMAX Results

```
=====
Dep. Variable:          y_true      No. Observations:          272
Model:                SARIMAX(1, 0, 0)  Log Likelihood          -2440.175
Date:                 Tue, 15 Sep 2020  AIC                   4884.350
Time:                 01:25:59         BIC                   4891.554
Sample:               0              HQIC                   4887.243
                        - 272
Covariance Type:      opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.9738	0.014	67.324	0.000	0.945	1.002
sigma2	3.878e+06	2.51e+05	15.424	0.000	3.39e+06	4.37e+06

```
=====
Ljung-Box (Q):          190.84    Jarque-Bera (JB):
31.08
Prob(Q):                0.00    Prob(JB):
0.00
Heteroskedasticity (H):  3.10    Skew:
-0.10
Prob(H) (two-sided):    0.00    Kurtosis:
4.65
=====
```

Warnings:

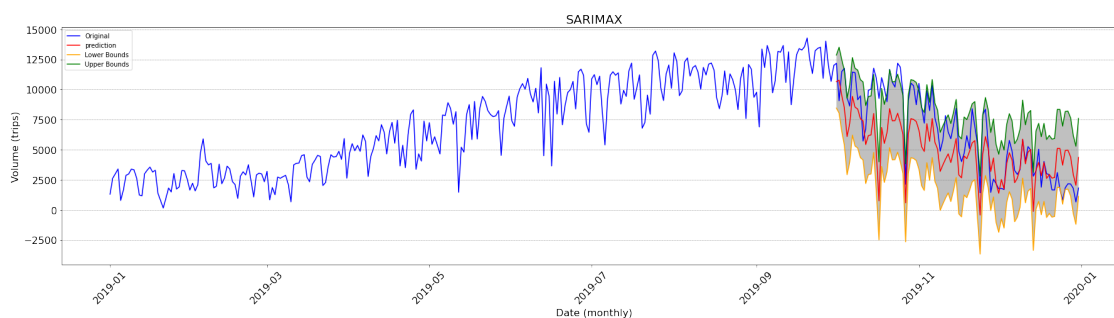
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

- From this result, although the MAPE is not very high, the prediction effect of the model is not ideal, especially the interval prediction effect.
- Next, try to add the exogenous variables to improve the model effect.

```
[373]: exog_lst_full = ['Temperature_Minimum', 'Temperature_Maximum',
↳ 'Temperature_Mean',
        'RelativeHumidity_Minimum', 'RelativeHumidity_Maximum',
↳ 'RelativeHumidity_Mean',
        'MeanSeaLevelPressure_Minimum',
↳ 'MeanSeaLevelPressure_Maximum', 'MeanSeaLevelPressure_Mean',
        'PrecipitationTotal_Summation', 'SnowfallAmount_Summation',
↳ 'CloudCoverTotal_Mean',
        'SunshineDuration_Summation', 'Evapotranspiration_Summation',
        'PBLHeight_Minimum', 'PBLHeight_Maximum', 'PBLHeight_Mean',
        'WindGust_Minimum', 'WindGust_Maximum', 'WindGust_Mean',
        'WindSpeed_Minimum', 'WindSpeed_Maximum', 'WindSpeed_Mean',
        'WindDirection_Dominant_None',
        'GeopotentialHeight_Minimum', 'GeopotentialHeight_Maximum',
↳ 'GeopotentialHeight_Mean',
        'day', 'dayofweek', 'is_workday', 'is_holiday']

result_df, model = sarimax_forecast(df=dataset, exog_lst=exog_lst_full,
                                   start_forecast_date=FORCAST_START_DATE,
                                   end_forecast_date=FORCAST_END_DATE)

arima_plot(result_df, title='SARIMAX')
model_eval(result_df['y_true'], result_df['y_pred'],
            start_date=FORCAST_START_DATE, end_date=FORCAST_END_DATE)
```



MAPE: 0.4362222118346819

MAD: 2065.216801665083

- After introducing exogenous variables to the model, the MAPE increases, the MAD decreases, and the model's predictions become more realistic.
- The increase of MAPE is partially because the low MAPE of the previous model is caused by

the generally predicted value being lower than the true value. So the higher MAPE here is relatively acceptable.

- In other words, adding exogenous variables does improve the model effect.

```
[358]: print(model.mle_retvals)
```

```
{'fopt': 8.373486516255104, 'gopt': array([ 2.44869014e-05,  1.27897692e-08,
4.23225900e-05, -6.89389879e-05,
      1.20093979e-04, -1.16954979e-04, -3.89289135e-04, -4.59752593e-04,
-4.10761559e-04,  4.60799398e-05, -3.85281140e-05,  1.51807633e-04,
-7.76026354e-04, -7.05451697e-05, -1.64280500e-04,  4.03786160e-04,
-7.48590523e-05, -1.44950718e-05, -7.66455344e-05, -7.12274684e-05,
-4.53752591e-05, -5.00328667e-05, -9.46016598e-06, -1.02211217e-04,
  3.30533645e-04, -8.31789748e-05,  1.60351554e-04, -2.93479019e-05,
  2.07300843e-07,  3.51732866e-05, -2.29828601e-05, -9.84173809e-04,
  7.32594431e-05]), 'fcalls': 15028, 'warnflag': 1, 'converged': False,
'iterations': 393}
```

- Lots of irrelevant exogenous variables cause the model fail to converge.
- Next, only useful variables are retained to make the model converge while maintaining a good model performance.

```
[359]: print(model.summary())
```

```

                                SARIMAX Results
=====
Dep. Variable:                  y_true      No. Observations:                  272
Model:                        SARIMAX(1, 0, 0)  Log Likelihood                  -2277.588
Date:                        Tue, 15 Sep 2020    AIC                           4621.177
Time:                        01:26:14          BIC                           4740.047
Sample:                      0                HQIC                       4668.904
                                - 272
Covariance Type:              opg
=====
=====
                                coef      std err          z      P>|z|
-----
[0.025      0.975]
-----
Temperature_Minimum           -181.8214     82.227     -2.211     0.027
-342.984      -20.659
Temperature_Maximum           -162.4794     95.886     -1.694     0.090
-350.413       25.455
Temperature_Mean              576.1300    163.945     3.514     0.000
254.803      897.457
RelativeHumidity_Minimum       -0.5622     15.748     -0.036     0.972
-31.427       30.303
RelativeHumidity_Maximum       11.9421     13.344     0.895     0.371
-14.211       38.096

```


RelativeHumidity_Mean	4.6847	24.426	0.192	0.848
-43.189 52.558				
MeanSeaLevelPressure_Minimum	-1106.9113	630.640	-1.755	0.079
-2342.943 129.120				
MeanSeaLevelPressure_Maximum	-883.0381	474.349	-1.862	0.063
-1812.745 46.669				
MeanSeaLevelPressure_Mean	1991.2926	1013.414	1.965	0.049
5.038 3977.547				
PrecipitationTotal_Summation	-98.6419	21.428	-4.603	0.000
-140.640 -56.643				
SnowfallAmount_Summation	27.0382	116.199	0.233	0.816
-200.708 254.784				
CloudCoverTotal_Mean	4.6729	10.424	0.448	0.654
-15.758 25.104				
SunshineDuration_Summation	1.9099	1.043	1.832	0.067
-0.134 3.954				
Evapotranspiration_Summation	150.6438	122.095	1.234	0.217
-88.657 389.945				
PBLHeight_Minimum	0.0997	0.891	0.112	0.911
-1.647 1.847				
PBLHeight_Maximum	0.0847	0.328	0.258	0.796
-0.558 0.727				
PBLHeight_Mean	0.8520	0.936	0.910	0.363
-0.983 2.687				
WindGust_Minimum	-1.5305	28.008	-0.055	0.956
-56.425 53.364				
WindGust_Maximum	-30.1916	22.140	-1.364	0.173
-73.586 13.203				
WindGust_Mean	16.5477	67.712	0.244	0.807
-116.165 149.260				
WindSpeed_Minimum	-66.8677	52.367	-1.277	0.202
-169.505 35.769				
WindSpeed_Maximum	-85.6664	54.653	-1.567	0.117
-192.784 21.451				
WindSpeed_Mean	115.7167	163.257	0.709	0.478
-204.260 435.694				
WindDirection Dominant_None	-1.9003	0.994	-1.911	0.056
-3.849 0.049				
GeopotentialHeight_Minimum	139.3753	77.181	1.806	0.071
-11.896 290.647				
GeopotentialHeight_Maximum	109.2087	58.610	1.863	0.062
-5.665 224.083				
GeopotentialHeight_Mean	-246.6524	124.262	-1.985	0.047
-490.202 -3.103				
day	-11.4208	20.442	-0.559	0.576
-51.486 28.644				
dayofweek	35.0937	64.963	0.540	0.589
-92.231 162.419				

```

is_workday          2015.2252    276.078    7.299    0.000
1474.122    2556.329
is_holiday          -1709.3151    325.055    -5.259    0.000
-2346.411    -1072.219
ar.L1               0.7322    0.058    12.655    0.000
0.619    0.846
sigma2             1.269e+06    1.12e+05    11.329    0.000
1.05e+06    1.49e+06
=====
===
Ljung-Box (Q):          118.39    Jarque-Bera (JB):
37.77
Prob(Q):                0.00    Prob(JB):
0.00
Heteroskedasticity (H):    1.92    Skew:
-0.30
Prob(H) (two-sided):      0.00    Kurtosis:
4.72
=====
===

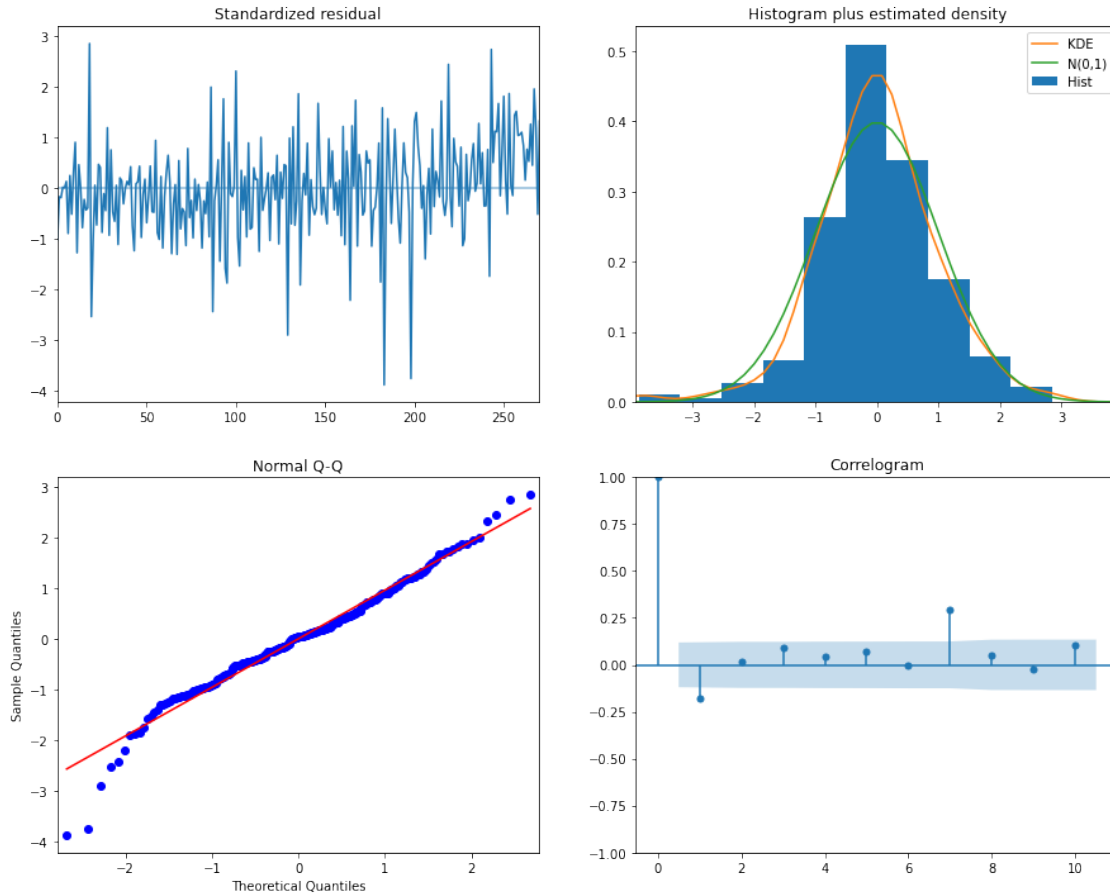
Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).

```

```

[360]: model.plot_diagnostics(figsize=(15, 12))
plt.show()

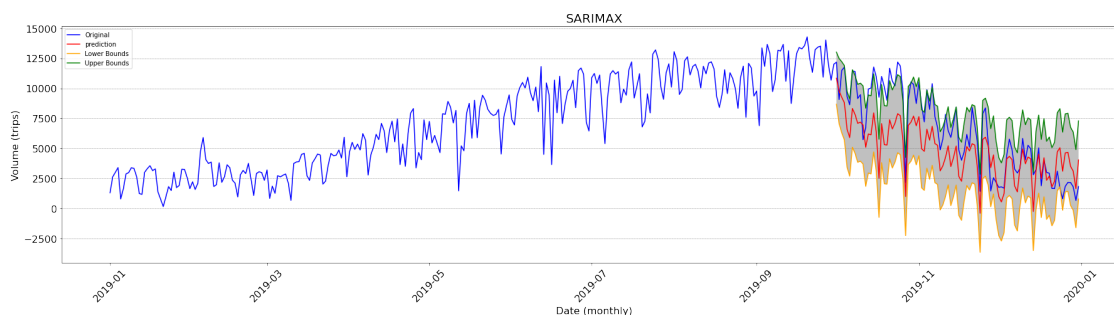
```



```
[376]: exog_lst_select = ['Temperature_Mean', 'MeanSeaLevelPressure_Mean',
    ↪ 'PrecipitationTotal_Summation',
    'Evapotranspiration_Summation', 'WindSpeed_Mean',
    'is_workday', 'is_holiday']

result_df, model = sarimax_forecast(df=dataset, exog_lst=exog_lst_select,
    start_forecast_date=FORCAST_START_DATE,
    end_forecast_date=FORCAST_END_DATE)

arima_plot(result_df, title='SARIMAX')
model_eval(result_df['y_true'], result_df['y_pred'],
    start_date=FORCAST_START_DATE, end_date=FORCAST_END_DATE)
```



MAPE: 0.4276895150879794

MAD: 2127.852603554388

```
[377]: print(model.mle_retvals)
```

```
{'fopt': 8.415744474999084, 'gopt': array([-7.47313322e-07,  9.78577219e-06,
 4.16910950e-07,  2.93631786e-07,
-2.29682939e-07,  2.24197550e-05, -1.96093808e-05,  1.32500055e-04,
-3.19579031e-05]), 'fcalls': 1190, 'warnflag': 0, 'converged': True,
'iterations': 101}
```

```
[379]: print(model.summary())
```

```

SARIMAX Results
=====
Dep. Variable:          y_true      No. Observations:          272
Model:                SARIMAX(1, 0, 0)  Log Likelihood          -2289.082
Date:                Tue, 15 Sep 2020    AIC                   4596.165
Time:                01:47:02           BIC                   4628.584
Sample:              0                 HQIC                  4609.182
                        - 272
Covariance Type:      opg
=====
=====

```

	coef	std err	z	P> z
[0.025 0.975]				

Temperature_Mean	207.1777	17.158	12.074	0.000
173.548 240.807				
MeanSeaLevelPressure_Mean	2.9217	0.430	6.788	0.000
2.078 3.765				
PrecipitationTotal_Summation	-104.3540	13.271	-7.863	0.000
-130.365 -78.343				
Evapotranspiration_Summation	617.4280	75.349	8.194	0.000
469.747 765.109				

```

WindSpeed_Mean          -90.0968    18.621    -4.838    0.000
-126.593    -53.601
is_workday              1850.6880    145.090    12.755    0.000
1566.316    2135.060
is_holiday              -1550.3419    217.792    -7.118    0.000
-1977.207    -1123.477
ar.L1                   0.7447     0.039    19.285    0.000
0.669     0.820
sigma2                  1.227e+06    7.91e+04    15.520    0.000
1.07e+06    1.38e+06
=====
===
Ljung-Box (Q):          101.39    Jarque-Bera (JB):
29.43
Prob(Q):                0.00    Prob(JB):
0.00
Heteroskedasticity (H):  3.74    Skew:
-0.14
Prob(H) (two-sided):    0.00    Kurtosis:
4.59
=====
===

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

- The model successfully converged and got better prediction results
- Useful exogenous variables are:
 - is today workday
 - is today holiday
 - Mean Temperature
 - Mean MeanSeaLevel Pressure
 - Precipitation Total Summation
 - Evapotranspiration Summation
 - Mean WindSpeed

6. [Optional] Focus on the trip data during 2019-01 and 2019-12 only. Now, the client of MSBA & Company wants to understand the key factors that explains the difference in the (average daily) ridership between different pairs of origin and destination. Please explore auxiliary data sets and discuss your findings. [3pts]